

COMPSYS 701 – Advanced Digital Design, Semester 1 2022

Designing Heterogeneous Multiprocessor System on Chip (HMPSoC) for Mixed-criticality Applications

(Research Based Course that is split into a Group Research Project (GRP) and Individual Research Project (IRP))

1. Introduction

The course centres on the topic of multi-core systems for mixed-critical embedded and real-time applications that clearly separate critical processing in terms of time predictability and safety from non-critical parts. The system uses physical isolation of these two parts of the execution platform and enables the two parts/subsystems communicate via standard bus and/or more general interconnection Network-on-Chip. Non-critical subsystem uses standard bus to connect communicating components, while critical subsystem is based on the use of time predictable NoC. Designing the system, which can be considered as Heterogeneous Multiprocessor System on Chip (HMPSoC), includes advanced digital design techniques when designing individual components and architecting them into a complex HMPSoC. Those components may include new general-purpose time predictable processor cores, application-specific processors and network interfaces, using some existing components in the form of IP blocks and their integration into HMPSoC. The course contains significant research as well as advanced hardware and software components.

2. Motivation and background

Increasing number of embedded and real-time systems demand computation requirements difficult to meet with standard available execution platforms. Those requirements go beyond capabilities of a single processor core and require multi-core platforms with large number of cores that have become a feasible and attractive alternative [1]. Communication requirements of programs that run on multi-core systems have naturally led towards various types of interconnect, e.g. standard buses, and networking of the cores in network-on-chip (NoC) based execution platforms [2]. While technology is not a bottleneck to design NoCs there are three issues that attract the attention of researchers: (1) what type of interconnect/NoC is suitable for certain types of execution platforms applications, (2) how to effectively use the cores to achieve target goals of the applications and (3) how to guarantee time-predictability of the execution platform when dealing with hard real-time and safety-critical parts of applications. These issues are very challenging and there are no definite answers. The cores of single type, which lead towards homogeneous multi-core systems, are not sufficient to satisfy requirements

of complex applications as the utilisation of cores and performance of the overall systems can be inadequate. The solutions could be achieved more effectively with the use of different types of cores, which together are able to achieve the goals of the target system in the best way in terms of targeted metrics such as response times, predictability, and energy consumption. Such Systems-on-Chip are referred to be heterogeneous as they comprise different types of cores [1]. In addition, employing heterogeneous MPSoCs helps optimising the power/energy consumption of the system. One approach that chip designers can provide is an execution platform and then the design of the embedded system applications using primarily software tools, which are not well developed for such platforms. The other approach can be to start the application development in a given tool and then synthesise the customised platform capable of executing the application by using certain off-the-shelf components (IP-Intellectual Property), but also allowing incorporating design specific components, and connecting them using a sort of standard interconnect, e.g. bus or/and NoC. In other words, the platform can be customised for specific application by using processing and storage components that need to be interconnected and enable transfer of data each to the other. This way, hardware/software partitioning of the application and co-design would be also supported. FPGAs are the best current practical technology that allow experiments with this design approach.

In this project we are working on the development of an execution platform for mixed-criticality systems that is heterogeneous in nature by incorporating different types of processor cores and application-specific components (in the form of IPs) that are interconnected using a combination of a standard bus and NoC. As the starting point we assume that we target real-time and safety critical applications which require guarantees in terms of functional properties as well as timing constraints for at least certain parts of the applications. We explore the options of integrating, customising and synthesising a target platform using different components such as our own simple time-predictable RISC processor called ReCOP, standard Intel's (Altera's) general purpose processor Nios II as well as our custom-made application-specific cores that accelerate the execution of signal processing or other algorithms. A general architecture of HMPSoC is illustrated in Figure 1.

In ideal case, our embedded and real-time system applications are developed in a new system-level language SystemGALS [3] that allows the use of a natural designer-specified parallelism and can exploit capabilities of multi-core execution platforms, including application-specific processors, and uses C as its host language. Although all tools for SystemGALS compilation and scheduling of executable code have not been developed yet, the paradigm of SystemGALS can be used to manually develop applications and generate partial executables of the overall systems as it was shown over years with its predecessor SystemJ GALS language [4] with fully automated tools. SystemGALS is a powerful concurrent programming and design language based on formal Globally Asynchronous Locally Synchronous (GALS) model of computation (MoC) that allows the specification and composition of very complex concurrent systems that target embedded and real-time applications [3, 4]. The language naturally supports explicit partitioning of the designed system on two types of computational units that correspond to two types of processing in complex embedded real-time systems:

- (1) control-driven processing based on GALS MoC that incorporates synchronous-reactive concurrent MoC for describing complex concurrent control flows and globally asynchronous concurrent MoC for composition of synchronous reactive components, and

- (2) data-driven processing where the computations require processing of data that arrives most often in streams and with high data rates. Additional motivation for using SystemGALS for the application design is that a designer can use it in different types of applications, from embedded non-real time systems to hard real-time and safety critical systems and the systems that are combination of both, we refer to them as mixed-criticality systems.

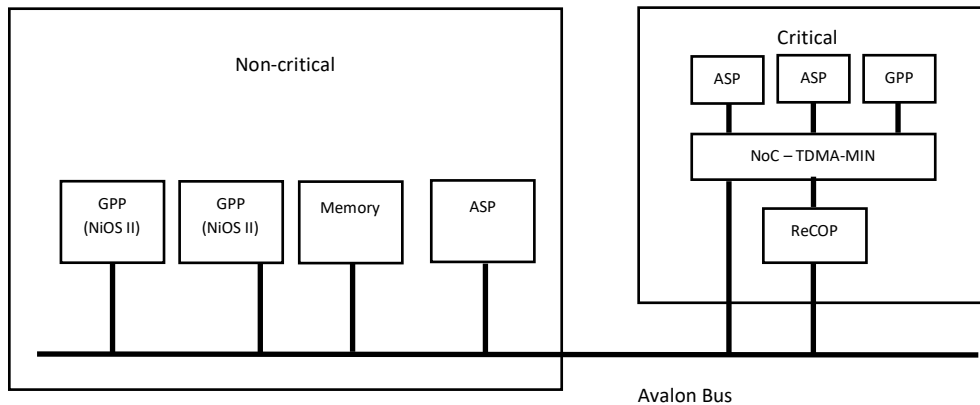


Figure 1 - General architecture of HMPSoC

3. HMPSoC – Heterogeneous Multiprocessor SoC

The overarching goal of the project is to investigate the options and develop a prototype of the HMPSoC, which allows system designers to combine programming and hardware solutions and hardware/software partitioning and design. HMPSoC uses two major interconnect mechanisms (see Figure 1), Avalon bus (Intel/Altera) that allows interconnection of various components to a single bus, and network-on-chip (NoC) developed specifically for time predictability necessary in critical part of the system [5] illustrated in Figure 2. Time-predictability is a pre-condition for static timing analysis of the execution time of the program or hardware component and estimation of worst-case execution times (WCET) and more general response times of the designed system. Here under “program” we generalise design specifications that are implemented in combination (composition) of software and hardware.

The HMPSoC’s critical part comprises different types of cores connected/networked using Time Division Multiple Access (TDMA) Multistage Interconnection Network (MIN), thus called TDMA-MIN [5], developed within Embedded Systems Research Group. A more detailed diagram that presents the TDMA-MIN interconnect fabric is shown in Figure 2.

TDMA-MIN allows connection of any nodes (processor cores, for example) via $N \times N$ switch-type interconnect, where N is maximal number of connected nodes. TDMA-MIN can be synthesised for any $N=2^k$ ($k=1, 2, 3, \dots$). In case of HMPSoC, it comprises at least one ReCOP processor core and the remaining cores are typically ASPs, but also can be general purpose processors (GPP) such as a Nios II processor core. The NoC provides data transfers between pairs of nodes without collision [6] and guarantees bounded latency (in terms of clock cycles) between any two nodes in NoC.

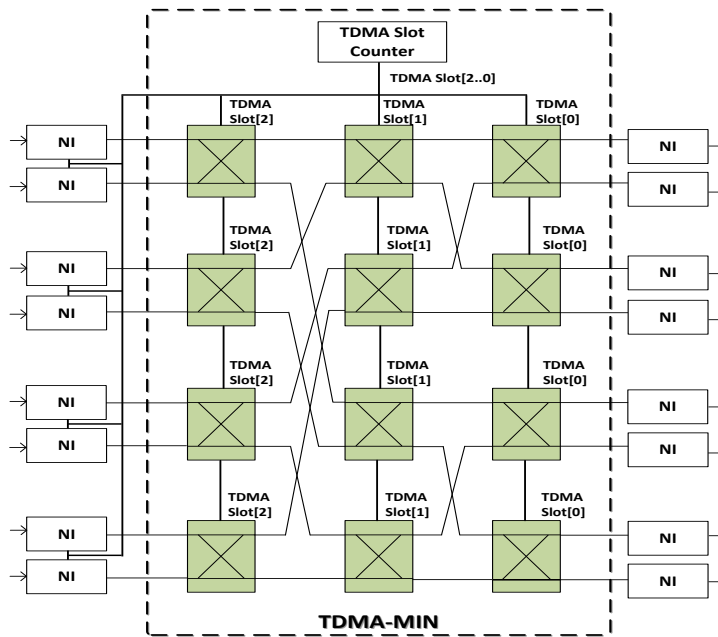


Figure 2 - TDMA-MIN Interconnect with the network interfaces (8x8 nodes)

4. Major goals (milestones) and constraints

The major goals of the GRP and IRP projects are:

1. Full understanding of the overall architecture of HMPSoC and its aims to enable design and implementation of instances of architecture customised to specific application or application domain. Understanding System-on-Chip (SoC) and Network-on-Chip (NoC) concepts generally and of TDMA-MIN NoC, specifically; understanding of all interconnects used to make the SoC (TDMA-MIN and Avalon-bus), and how interfaces of individual nodes are made to provide connectivity.
2. Ability to critically read, analyse and assess existing research and other literature including existing provided designs.
3. Ability to design fully operational ReCOP RISC processor with necessary customisation including the assembler that supports the addition of new instructions.
4. Design an application specific processor (ASP) that can be interfaced with TDMA-MIN NoC and/or Avalon bus; functionality of the ASP will be specified separately within Individual Research Project (IRP); the ASP is in the domain of performing digital signal processing algorithms in custom digital hardware and implemented in FPGA.
5. Integration of components into SoC using Intel FPGA Platform Designer.
6. Use of VHDL to design the required components of the system at RTL level and to simulate before synthesis using ModelSim.
7. Prototyping with Intel/Altera FPGA on DE1-SoC and/or DE2-115 prototyping board.
8. Testing parts of the system with software developed using SystemGALS style specifications; those specifications will be manually mapped on software in C and ReCOP assembly language.
9. Working in a team with understanding own and other members' responsibilities and roles and respecting the contributions of all members of the team.

In order to be manageable and achievable, the project will assume the following design constraints as a starting point:

1. A single type of new general purpose RISC processor called Reactive and Concurrency-Processor (ReCOP) [7, 8, 9] that allows easy customisation to satisfy the requirements of the applications by simple internal changes, and Nios II as given mature general purpose processor (GPP). While Nios II processor(s) are primarily used in implementation of non-critical part of the system, they are also used as support environment for configuration of critical part.
2. Different types of Application-Specific Processor (ASP) cores, some of them are specified further in this brief and others in separate IRP brief.
3. Network-on-Chip based on TDMA-MIN interconnect [5, 6].

Due to limited project duration, our goal is to develop an experimental platform that has constrained capacity. For the critical part, it comprises at least one ReCOP, 3 ASPs and provides 8 ports in the TDMA-MIN NoC, which allows connecting up to 8 cores, ReCOPs and ASPs. One of the ports can be used to extend the connectivity to Avalon bus and non-critical part of the system.

As one of the sub-goals is to master full, yet simple, processor design; the provided ReCOP core initial specification, given in the form of high-level programming model of the processor that includes Instruction Set Architecture (ISA) needs to be expanded to full design and implemented as register transfer level (RTL) design.

Project will be conducted through a number of phases and tasks in order to make it easier to manage and reduce the complexity of design process. After the analysis and understanding of the overall goals, a careful plan and alignment with each phase is the key for the successful completion of the tasks. Besides the team effort, which constitutes 60% of the marks, the project contains an individual component of exploring and designing an ASP, which constitutes 40% of the final grade of each student. The project is done in teams with four members. All issues in the team function, operation and collaboration have to be indicated immediately to the academic staff and course coordinator by either the team or individual members of the team.

5. Background on processor design

In order to prepare for processor design, within the lectures we are going to present full design of an early pioneering reactive microprocessor called ReMIC, including rationale for its design, description of its programming model, ISA, datapath and control unit, as well as potential for user and application-specific customisation. First, it will include the design of non-reactive core of the processor [10], followed by the design of customisation for reactive systems [11] in multi-cycle version and then pipelined version will be introduced.

6. Application domain for ASP

We envisage as a target any application which requires certain level of

- (1) Reactivity on external inputs and events (such as pressing of push buttons, sudden changes on external inputs, which indicate the need for change of control flow within the application and internally generated events that mean sudden change of signal values detected during system operation).

- (2) Processing of streams of data, which come to the system (sourced from analogue signal and ADC converter) presented as time series of samples, processing of those samples using signal processing algorithms as specified below or cascade of those algorithms, and conversion of the digital results into the output stream/time series ready to be converted into analogue form (via DAC converter).
- (3) Support and non-critical functions that will be used in system operation that execute on Nios II processor (such as configuration of critical part of the platform, preparation of data, change of code of critical part, debugging, customisation of instruction set, communication with the designer team etc).

For the purposes of this course you will design a simple tool for signal analysis (much like an oscilloscope), consisting of a single ReCOP processor, a set of Application Specific Processors (ASPs) for data processing, and input and output ASPs emulating data interaction with the real world, all interacting via a TDMA-MIN NoC. The real time processor will handle control inputs (switches and/or buttons) and state outputs (LEDs) and configure the ASPs and NoC to alter the behaviour of the system at runtime. Configuration and data will be transferred using standalone 32-bit operations to avoid the need for buffering data across the NoC.

In individual research project (IRP) each student will have to implement an ASP with a set of features described in IRP brief. ASPs will operate on streaming samples of data, perform computations over an internal buffer implemented in local memory, and output data as appropriate. Examples of typical computations on the data streams are:

- Passthrough of data
- Linear filters
- Peak minimum / maximum detection in the time series
- Average over a moving/sliding window
- Finite sum-of-products calculation for two data series
- etc

An Input ASP (emulating an ADC) can be made with a set of virtual channels and sampling rates, as well as an Output ASP that provides an analogue output from the prototyping board; they will be also referred to as IO-ASPs. ASPs that perform data processing (computations) will be referred to as DP-ASP. Sometimes, if IO-ASPs have more specific functionality, for example emulating ADC and DAC operation, they can be referred to as ADC-ASP and DAC-ASP, respectively.

For the initial phases you may use a customised state machine ASP (instead of ReCOP) to send configuration messages over the TDMA-MIN interface and configure other ASPs; this will allow teams to work on the NoC and ASPs in parallel with the development of ReCOP.

The functional application will be expected to take/read a configuration values from switches on the development board to:

- Select the input channel and sampling rate
- Enable a combination of each of the ASP provided filters for each channel:
 - Direct passthrough
 - Linear filtering

- Averaging
- Peak Detection

Data formats and type of expected arithmetic will be determined separately and kept at fixed point or integer arithmetic only. There will be no need to design or use any dividers. Division operation will be limited to the numbers that are 2^k , where k is limited to 0, 1, 2 and 3.

- As an example, to configure linear filter mode the ReCOP would send a DAC-ASP configure command to enable the appropriate DAC output, then a Data Processing ASP (DP-ASP) command to configure it to execute a linear filter over the provided data with the DAC as an output destination, then finally an ADC-ASP configuration command to select the ADC channel, sample rate, and set the output destination to the DP-ASP. At this point the ADC-ASP will begin sending data to the DP-ASP, then the DP-ASP will perform the linear filtering function and (when data is available) this will send the data to the DAC-ASP. Similarly, you can create different datapaths that process the data stream in cascade that comprises different processing blocks (DP-ASPs).

The current system state should be reflected on the development board LEDs, and outputs will be observed from the Output ASP (O-ASP) using an oscilloscope.

7. Timeline-deliverables-assessment

Major tasks and checkpoints and tentative dynamics

Week	1	2	3	4	5	6	7	B1	B2	8	9	10	11	12
Task/milestone														
Conceptual system design/				FB										
Nios II/CI and Avalon bus					GRP Ass-1 10%									
ReCOP										GRP (ReCOP)				
IRP (DP-ASP with NoC Interface/NI)							IRP Phase 1 Ass-2 10%				IRP Phase 2 (30%)			
Integration Critical Part												GRP Phase 1 30%		
HMPSoC integration														GRP Phase 2 20%

FB- Feedback

- Assignment – 1 is part of GRP preparation (GRP Phase 0), which can be done by a group of two students.
- Assignment – 2 is part of IRP (as IRP Phase 1), which should be done individually.
- IRP Phase 2- individual work as specified in the IRP project brief
- GRP Phase 1- Integration of critical part of HMPSoC which includes development of ReCOP microprocessor as an intermediate milestone
- GRP Phase 2 – Integration of Non-critical and Critical Part of HMPSoC

8. Resources and tools

- DE1-SoC Terasic FPGA Development Kit for Altera/Intel Cyclone V based systems
(An alternative can be DE2-115 FPGA Development Kit with Cyclone IV FPGA device if DE1-SoC is not available)

- Intel Quartus Development Environment including Platform Designer
- ModelSim
- ReCOP assembler
- Other resources will be provided via Canvas

9. References and readings

1. Singh, A.K., Shafique, M., Kumar, A. and Henkel, J., 2013, May. Mapping on multi/many-core systems: survey of current and emerging trends. In *Proceedings of the 50th Annual Design Automation Conference* (p. 1). ACM.
2. Benini, L. and De Micheli, G. 2002. Networks on chips: a new SoC paradigm. *Computer* 35, 70-78.
3. Salcic, Z., Park, H., Biglari-Abhari, M., and Teich, J., 2019, SystemGALS – A language for the design of GALS software systems, Embedded Systems Research Group, University of Auckland, Internal document, Embedded Systems Research Group, available to C701 class
4. Malik, A., Salcic, Z., Roop, P.S. and Girault, A. 2010. SystemJ: A GALS language for system level design. *Computer Languages, Systems & Structures* 36, 317-344
5. Salcic, Z., Nadeem, M. and Striebing, B, 2016, A Time Predictable Heterogeneous Multicore Processor for Hard Real-time GALS Programs. *ARCS 2016*
6. Salcic Z, Park H, Teich J, Malik A, Nadeem M , 2017, NoC-HMP: A Heterogeneous Multicore Processor for Embedded Systems Designed in SystemJ, *ACM Transactions on Design Automation of Embedded Systems*, 2017, Volume 22 Issue 4
7. Salcic, Z., Nadeem, M., Park, H. and Teich, J, 2016, Optimizing Latencies and Customizing NoC of Time-Predictable Heterogeneous Multi-Core Processor. In *2016 IEEE 10th International Symposium on Embedded Multicore/Many-core Systems-on-Chip*, Lyon
8. Salcic, Z & Lorigan, H, 2020, ReCOP – A Processor Core for Control-dominated Reactive Applications, University of Auckland, Department of Electrical and Computer Engineering, Embedded Systems Research Group, a modified version of 2015 document, document available to the C701 class
9. Salcic, Z. and Malik, A. 2013. GALS-HMP: A heterogeneous multiprocessor for embedded applications. *ACM Transactions on Embedded Computing Systems (TECS)* 12, 58.
10. Hui, D. and Salcic, Z., 2004, MiCORE- A Customisable Microprocessor Core, University of Auckland, Department of Electrical and Computer Engineering, Embedded Systems Research Group Internal Report
11. Hui, D. and Salcic, Z., 2004, ReMIC- A Customisable Reactive Microprocessor Core, University of Auckland, Department of Electrical and Computer Engineering, Embedded Systems Research Group Internal Report